

Learning Paths

Training Courses

Certification

[CCP: Data Scientist](#)

Hadoop Developer CCDH

Hadoop Admin CCAH

HBase Specialist CCSHB

Online Resources

Private Training

Training Partners

1. Exploring the Data Set

Before you begin any work, you should familiarize yourself with the data, both its condition and the elements you want to extract in the context of the questions you want to answer.

Step 1. Profile the data

Start with some general data profiling.

1. Determine how much data you're working with.

Open a terminal (by right-clicking on the desktop and selecting *Open in Terminal*) and execute the following commands:

```
$ cd ~/data
$ du -sh .
201M .
```

The `du` command reports the current directory's disk usage. List the contents of the directory and its subdirectories (recursively) with the following command:

```
$ ls -Rlh
total 0
drwxr-xr-x 12 training staff 408B Oct 18 12:53 heckle
drwxr-xr-x 12 training staff 408B Oct 18 12:53 jeckle

./heckle:
total 103M
-rw-r--r-- 1 training staff 18M Oct 18 12:48 web.log
-rw-r--r-- 1 training staff 17M Oct 18 12:48 web.log.1
-rw-r--r-- 1 training staff 253K Oct 18 12:48 web.log.2
-rw-r--r-- 1 training staff 14M Oct 18 12:48 web.log.3
-rw-r--r-- 1 training staff 11M Oct 18 12:48 web.log.4
-rw-r--r-- 1 training staff 2.1M Oct 18 12:48 web.log.5
-rw-r--r-- 1 training staff 655K Oct 18 12:48 web.log.6
-rw-r--r-- 1 training staff 11M Oct 18 12:50 web.log.7
-rw-r--r-- 1 training staff 15M Oct 18 12:50 web.log.8
-rw-r--r-- 1 training staff 14M Oct 18 12:51 web.log.9

./jeckle:
total 99M
-rw-r--r-- 1 training staff 17M Oct 18 12:53 web.log
-rw-r--r-- 1 training staff 16M Oct 18 12:53 web.log.1
-rw-r--r-- 1 training staff 3.5K Oct 18 12:53 web.log.2
-rw-r--r-- 1 training staff 14M Oct 18 12:53 web.log.3
-rw-r--r-- 1 training staff 9.1M Oct 18 12:53 web.log.4
-rw-r--r-- 1 training staff 2.0M Oct 18 12:53 web.log.5
-rw-r--r-- 1 training staff 504K Oct 18 12:53 web.log.6
-rw-r--r-- 1 training staff 11M Oct 18 12:53 web.log.7
-rw-r--r-- 1 training staff 14M Oct 18 12:53 web.log.8
-rw-r--r-- 1 training staff 14M Oct 18 12:53 web.log.9
```

2. Determine the number of files and lines

```
$ find . -type f -print | wc -l
20
```

```
$ find . -type f | xargs -n 1 wc -l | awk '{sum+=$1} END {print sum}'
612873
```

The `find` command outputs the names of files that match a set of filters, and the `wc -l` command counts lines. The last command uses `find` to list all files (not directories), `xargs` and `wc -l` to count the lines in those files, and `awk` to sum up those counts.

From these results, you can see that you're working with 200MB of data spread across 20 files containing over 60k total log lines.

3. Inspect the log data

Take a look at what the log data looks like. The command below looks at the first five lines.

```
$ head -5 heckle/web.log
{"auth": "15a63c4:e66189ba", "createdAt": "2013-05-12T00:00:01-08:00", "payload": {"itemId": "15607",
"marker": 240}, "refId": "47c7e2f6", "sessionId": "82ada851-0b3c-4e9d-b8cf-0f0a2ebed278", "type":
"Play", "user": 22700996, "userAgent": "Mozilla/4.0 (compatible; MSIE 7.0; Windows NT 5.1; .NET CLR
1.1.4322; .NET CLR 2.0.50727; InfoPath.1)"}
{"auth": "1547142:7d3d41c7", "createdAt": "2013-05-12T00:00:03-08:00", "payload": {"itemId": "6210",
"marker": 3420}, "refId": "141ac867", "sessionId": "d95bc727-033f-4f62-831a-2f8d6740a364", "type":
"Play", "user": 22311234, "userAgent": "Mozilla/5.0 (Windows NT 6.0) AppleWebKit/535.7 (KHTML, like
Gecko) Chrome/16.0.912.75 Safari/535.7"}
{"auth": "30af4f8:2527ff80", "createdAt": "2013-05-12T00:00:09-08:00", "payload": {"itemId": "32009",
"marker": 2760}, "refId": "fdec4481", "sessionId": "673ee60a-0aa2-4eac-a6fb-8a68d053dbf3", "type":
"Play", "user": 51049720, "userAgent": "Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/535.19 (KHTML,
like Gecko) Chrome/18.0.1025.142 Safari/535.19"}
{"auth": "6f691c:455e17cb", "createdAt": "2013-05-12T00:00:10-08:00", "payload": {"itemId": "7347",
"marker": 1059}, "refId": "4b5021f4", "sessionId": "2d3aef1d-ec8d-4053-8c40-e8579e547745", "type":
"Play", "user": 7301404, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 5.1; Trident/4.0;
WDL6.1.1; .NET CLR 1.1.4322; .NET CLR 2.0.50727)"}
{"auth": "1208d4c:279737f7", "createdAt": "2013-05-12T00:00:11-08:00", "payload": {"itemId":
"3702e4", "marker": 780}, "refId": "7586e549", "sessionId": "d4a244cb-d502-4c94-a80d-3d26ca54a449",
"type": "Play", "user": 18910540, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1;
Trident/4.0; GTB7.2; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center
PC 6.0; InfoPath.2)"}`
```

The `head -5` command prints the first five lines of the file or input.

You should recognize that the logs are in JSON format. They include `user`, `timestamp`, `user agent`, and `sessionId` for all entries in the sample. The logs also include an `itemId` in the `payload` field for all entries. As all entries in the sample are for "Play" events, you can deduce that the `itemId` indicates what was played.

Step 2. Load the data into HDFS

Because you have a large amount of data, you'll want to use Hadoop to parallelize your exploration operations.

1. Load using `-copyFromLocal`

Load all the data into HDFS with the following commands:

```
$ cd ..
$ pwd
/home/cloudera
$ hadoop fs -ls
$ hadoop fs -copyFromLocal data
```

The `hadoop fs -ls` command shows the files in your home directory in HDFS. There are currently none, so no output is displayed. The `hadoop fs -copyFromLocal` command copies the specified file or directory contents — in this case, the directory `data` — into HDFS. If only one argument is provided, the file or directory will be copied into the user's home directory in HDFS. When the command completes, the data will be available in HDFS under the `/user/cloudera/data` path.

2. Ensure the data loaded

Enter the following commands:

```
$ hadoop fs -ls
Found 1 item
drwxr-xr-x - cloudera cloudera      0 2014-02-13 15:19 data
$ hadoop fs -ls data
Found 2 items
drwxr-xr-x - cloudera cloudera      0 2014-02-13 15:19 data/heckle
drwxr-xr-x - cloudera cloudera      0 2014-02-13 15:19 data/jeckle
```

Using `hadoop fs -ls` you can see that the data has been uploaded, and the data directory contains the correct folders.

Step 3. Determine event types

Now, determine what other events the data reports other than play events. To do that, you need to write a simple Hadoop streaming script. You have to use a script because Hadoop streaming doesn't handle nested quotes in a command line argument very well.

1. Write a Hadoop streaming script

Create a file in the current directory titled `grep_field.sh` with the following contents:

```
#!/bin/bash
grep -Eo "\"$1\": [^,]+\" | cut -d: -f2- | tr -d ' ' '
```

The `grep` command filters its input according to the given patterns and outputs only matching lines. The `-E` option turns on extended regular expressions, and the `-o` option causes `grep` to only output the portion of the input that matches the pattern.

2. Run the script

Save the file in the current directory (`/home/cloudera`) and run it with the following command:

```
$ hadoop jar $STREAMING -D stream.non.zero.exit.is.failure=false -input data/heckle/ -input
data/jeckle/ -output types -mapper "grep_field.sh type" -file grep_field.sh -reducer "uniq -c"
```

The `-mapper` and `-reducer` options tell Hadoop what to execute as the map and reduce tasks, respectively. The `-file` command tells Hadoop to make sure the given file is available on all nodes that will execute a task. The `-input` and `-output` options tell Hadoop where to find the data and where to put the results. The `-D` argument is required because `grep` returns `1` if it doesn't find anything, so you have to tell Hadoop to ignore the return code when deciding if the task failed. (A return code of `1` is otherwise taken to indicate task failure.)

After the command completes, look at the results in your terminal window:

```
$ hadoop fs -cat types/part\*
177 "type": "Account"
5091 "type": "AddToQueue"
3062 "type": "Advance"
5425 "type": "Home"
19617 "type": "Hover"
274 "type": "ItemPage"
1057 "type": "Login"
1018 "type": "Logout"
4424 "type": "Pause"
558568 "type": "Play"
164 "type": "Position"
1313 "type": "Queue"
652 "type": "Rate"
1344 "type": "Recommendations"
1774 "type": "Resume"
1328 "type": "Search"
7178 "type": "Stop"
133 "type": "VerifyPassword"
274 "type": "WriteReview"
```

The `hadoop fs -cat` command outputs the contents of the named HDFS file.

You now see a list of events and the number of times they occur in the data. An overwhelming majority of events are "Play" events. The data also contains "Position", "Pause", "Advance", "Resume", and "Stop" events, which likely also relate to content playback. There are also a variety of other event types.

One question that immediately arises is whether all lines have a "type" field. To determine whether all lines have a type field, enter the following command:

```
$ hadoop fs -cat types/part-00000 | awk '{sum+= $1} END {print sum}'
612873
```

The result is the same result that the line count gave us, so you know that all lines in the log do indeed have a type field.

3. Review the events fields (a single example event of each type)

Next, look at some example lines from the data to see what these events look like. Because you'll be performing very basic operations, it will be faster and simpler to use UNIX commands in the Hadoop Shell directly instead of MapReduce. Run the following command:

```
$ hadoop fs -cat types/part\* | awk '{ print $2 }' | xargs -n 1 grep -Rhm 1 data -e
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:01:27-08:00", "payload": {"new": "kid",
"old": "adult", "subAction": "parentalControls"}, "refId": "752c4bc5", "sessionId": "81076d3e-ad42-
4567-a5dd-df3d4a0f3274", "type": "Account", "user": 52029279, "userAgent": "Mozilla/4.0 (compatible;
MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR
3.0.30729; .NET4.0C; .NET4.0E; .NET CLR 1.1.4322; InfoPath.3)"}
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:01:11-08:00", "payload": {"itemId": "7215"},
"refId": "752c4bc5", "sessionId": "81076d3e-ad42-4567-a5dd-df3d4a0f3274", "type": "AddToQueue",
"user": 52029279, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0;
SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; .NET4.0C; .NET4.0E; .NET CLR
1.1.4322; InfoPath.3)"}
{"created_at": "2013-05-08T08:04:10Z", "payload": {"item_id": "10965", "marker": 1685}, "session_id":
"0f0fe36d-359a-4d65-a865-a0d966e1a953", "type": "Advance", "user": 34534121, "user_agent":
"Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.0; Trident/4.0; SLCC1; .NET CLR 2.0.50727;
InfoPath.2; .NET CLR 1.1.4322; .NET CLR 3.0.30618; .NET CLR 3.5.30729; .NET4.0C; .NET4.0E; MS-RTC LM
8)"}
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:00:54-08:00", "payload": {"popular":
["32421", "10515", "7215", "30915", "37830"], "recent": ["17863e19"], "recommended": ["26841",
"12663e23", "5573", "29304", "4677"]}, "refId": "752c4bc5", "sessionId": "81076d3e-ad42-4567-a5dd-
df3d4a0f3274", "type": "Home", "user": 52029279, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0;
Windows NT 6.1; WOW64; Trident/4.0; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729;
.NET4.0C; .NET4.0E; .NET CLR 1.1.4322; InfoPath.3)"}
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:01:03-08:00", "payload": {"itemId": "7215"},
"refId": "752c4bc5", "sessionId": "81076d3e-ad42-4567-a5dd-df3d4a0f3274", "type": "Hover", "user":
52029279, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; SLCC2;
.NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; .NET4.0C; .NET4.0E; .NET CLR 1.1.4322;
InfoPath.3)"}
{"auth": "1208d4c:279737f7", "createdAt": "2013-05-12T00:51:22-08:00", "payload": {"itemId":
"3702e6"}, "refId": "7586e549", "sessionId": "d4a244cb-d502-4c94-a80d-3d26ca54a449", "type":
"ItemPage", "user": 18910540, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1;
Trident/4.0; GTB7.2; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center
PC 6.0; InfoPath.2)"}
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:00:47-08:00", "refId": "752c4bc5",
"sessionId": "81076d3e-ad42-4567-a5dd-df3d4a0f3274", "type": "Login", "user": 52029279, "userAgent":
"Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; SLCC2; .NET CLR 2.0.50727;
.NET CLR 3.5.30729; .NET CLR 3.0.30729; .NET4.0C; .NET4.0E; .NET CLR 1.1.4322; InfoPath.3)"}
{"auth": "1e7a371:396e2d82", "createdAt": "2013-05-12T00:26:13-08:00", "refId": "2b346725",
"sessionId": "c475f872-b28a-40eb-9f63-78584a3dbdb9", "type": "Logout", "user": 31957873, "userAgent":
"Mozilla/5.0 (Windows NT 6.1) AppleWebKit/535.19 (KHTML, like Gecko) Chrome/18.0.1025.33
Safari/535.19"}
{"auth": "6f691c:455e17cb", "createdAt": "2013-05-12T00:03:46-08:00", "payload": {"itemId": "7347",
"marker": 1275}, "refId": "4b5021f4", "sessionId": "2d3aef1d-ec8d-4053-8c40-e8579e547745", "type":
"Pause", "user": 7301404, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 5.1;
Trident/4.0; WDL6.1.1; .NET CLR 1.1.4322; .NET CLR 2.0.50727)"}
{"auth": "15a63c4:e66189ba", "createdAt": "2013-05-12T00:00:01-08:00", "payload": {"itemId": "15607",
"marker": 240}, "refId": "47c7e2f6", "sessionId": "82ada851-0b3c-4e9d-b8cf-0f0a2ebed278", "type":
"Play", "user": 22700996, "userAgent": "Mozilla/4.0 (compatible; MSIE 7.0; Windows NT 5.1; .NET CLR
```

```

1.1.4322; .NET CLR 2.0.50727; InfoPath.1)"}
{"created_at": "2013-05-08T08:01:57Z", "payload": {"item_id": "20435", "marker": 3103}, "session_id":
"0964d73f-f977-4673-961d-488b7c230aa5", "type": "Position", "user": 48418325, "user_agent":
"Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.9.2.9) Gecko/20110412 CentOS/3.6.9-2.el6.centos
Firefox/3.6.9"}
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:01:11-08:00", "payload": {"itemId": "7215"},
"refId": "752c4bc5", "sessionId": "81076d3e-ad42-4567-a5dd-df3d4a0f3274", "type": "AddToQueue",
"user": 52029279, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0;
SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; .NET4.0C; .NET4.0E; .NET CLR
1.1.4322; InfoPath.3)"}
{"auth": "4cbf0d:563049ae", "createdAt": "2013-05-12T00:01:07-08:00", "payload": {"itemId": "19474",
"rating": 5}, "refId": "28578fab", "sessionId": "90b8fa01-bcf4-44d7-b218-54b49b0b1064", "type":
"Rate", "user": 5029645, "userAgent": "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64;
Trident/4.0; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0;
InfoPath.3; MS-RTC LM 8; .NET4.0C; .NET4.0E)"}
{"auth": "1c177a5:77d90472", "createdAt": "2013-05-12T00:01:50-08:00", "payload": {"recs": ["11434",
"4697", "1342", "7579", "25865", "30393", "24789e45", "6324", "22407", "12276", "17161e1", "8276",
"18569", "16687", "1728e133", "13862", "26706", "30251", "32929", "4180", "6160", "31310", "10965",
"34721", "39155"]}, "refId": "4a7136d3", "sessionId": "c370df82-1dab-4c76-b0cc-ff0747d5b8b7", "type":
"Recommendations", "user": 29456293, "userAgent": "Mozilla/5.0 (Windows NT 6.1; WOW64)
AppleWebKit/536.11 (KHTML, like Gecko) Chrome/20.0.1132.43 Safari/536.11"}
{"created_at": "2013-05-08T08:02:01Z", "payload": {"item_id": "20435", "marker": 3103}, "session_id":
"0964d73f-f977-4673-961d-488b7c230aa5", "type": "Resume", "user": 48418325, "user_agent": "Mozilla/5.0
(X11; U; Linux i686; en-US; rv:1.9.2.9) Gecko/20110412 CentOS/3.6.9-2.el6.centos Firefox/3.6.9"}
{"auth": "2af23fc:789efe49", "createdAt": "2013-05-12T00:02:29-08:00", "payload": {"results":
["38759", "33586e1", "34199", "2955", "34465", "27184", "31020", "39834", "24668", "39164", "23999",
"16379e1", "18356", "16704", "35935", "21216", "26857e1", "19194", "2152", "32340"]}, "refId":
"2b584b36", "sessionId": "17868eda-5456-49ad-b6cd-aafcda16e592", "type": "Search", "user": 45032444,
"userAgent": "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_7_2) AppleWebKit/534.52.7 (KHTML, like
Gecko)"}
{"auth": "15a63c4:e66189ba", "createdAt": "2013-05-12T00:00:25-08:00", "payload": {"itemId": "15607",
"marker": 300}, "refId": "47c7e2f6", "sessionId": "82ada851-0b3c-4e9d-b8cf-0f0a2ebd278", "type":
"Stop", "user": 22700996, "userAgent": "Mozilla/4.0 (compatible; MSIE 7.0; Windows NT 5.1; .NET CLR
1.1.4322; .NET CLR 2.0.50727; InfoPath.1)"}
{"auth": "4497424:66fc2649", "createdAt": "2013-05-12T02:03:34-08:00", "refId": "7b172ae0",
"sessionId": "367feb29-f46d-49cd-9fde-537aa2ddfd6f", "type": "VerifyPassword", "user": 71922724,
"userAgent": "Mozilla/5.0 (Windows NT 5.1; rv:10.0.1) Gecko/20100101 Firefox/10.0.1"}
{"auth": "f6a6d2:4b2c55f0", "createdAt": "2013-05-12T00:02:39-08:00", "payload": {"itemId":
"37072e3", "length": 968, "rating": 3}, "refId": "6663e11a", "sessionId": "e6e02174-34d8-41f6-b723-
29d637e493d9", "type": "WriteReview", "user": 16164562, "userAgent": "Mozilla/4.0 (compatible; MSIE
8.0; Windows NT 5.1; Trident/4.0; .NET CLR 1.1.4322; .NET CLR 2.0.50727; .NET CLR 3.0.04506.30; .NET
CLR 3.0.04506.648; .NET CLR 3.0.4506.2152; .NET CLR 3.5.30729; InfoPath.2; MS-RTC LM 8)"}

```

This command uses `hadoop fs -cat` and `awk` to extract the set of unique types, and `xargs` and `grep` to find one occurrence of a line in the data containing each of the unique types.

The results return a single example event of each type. Browsing through the results, you should notice a common structure to the log lines. Every line has a type, user id, session id, creation timestamp, some kind of reference id, some kind of authorization code, user agent, and payload fields. The payload field appears to contain additional information about the event.

Step 4. Summarize the data

In this step, you want to determine whether all log lines share the common structure you discovered in the previous step. To do this, you'll need a MapReduce job that can handle JSON. These examples use Python as it has excellent JSON support and is rapidly becoming a favorite tool for data scientists.

1. Write a mapper

You will need a map task that can emit every field in the data. You can then sum up those occurrences in the reduce phase to get total counts. Create a file called `summary_map.py` with the following contents:

```

#!/usr/bin/python
import json
import sys
# Read all lines from stdin
for line in sys.stdin:

```

```
# Parse the JSON
data = json.loads(line)
# Emit every field
for field in data.keys():
    print field
```

2. Run the job

Enter the following command to run the job, again using `uniq -c` as the reducer which aggregates the fields and returns the counts to you.

```
$ hadoop jar $STREAMING -input data/heckle/ -input data/jeckle/ -output summary -mapper summary_map.py
-file summary_map.py -reducer "uniq -c"
packageJobJar: [../summary_map.py, /tmp/hadoop-training/hadoop-unjar5049687800382317379/] []
/var/folders/1l/xl67db9x2rs47q5fs4b255rr0000gp/T/streamjob7073892776336205044.jar tmpDir=null
13/10/23 11:11:49 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments.
Applications should implement Tool for the same.
13/10/23 11:11:49 INFO mapred.FileInputFormat: Total input paths to process : 20
13/10/23 11:11:49 INFO streaming.StreamJob: getLocalDirs(): [/tmp/hadoop-training/mapred/local]
13/10/23 11:11:49 INFO streaming.StreamJob: Running job: job_201310231802_0004
13/10/23 11:11:49 INFO streaming.StreamJob: To kill this job, run:
13/10/23 11:11:49 INFO streaming.StreamJob: /usr/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -
kill job_201310231802_0004
13/10/23 11:11:49 INFO streaming.StreamJob: Tracking URL: http://localhost:50030/jobdetails.jsp?
jobid=job_201310231802_0004
13/10/23 11:11:50 INFO streaming.StreamJob: map 0% reduce 0%
13/10/23 11:12:18 INFO streaming.StreamJob: map 100% reduce 100%
13/10/23 11:12:18 INFO streaming.StreamJob: To kill this job, run:
13/10/23 11:12:18 INFO streaming.StreamJob: /usr/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -
kill job_201310231802_0004
13/10/23 11:12:18 INFO streaming.StreamJob: Tracking URL: http://localhost:50030/jobdetails.jsp?
jobid=job_201310231802_0004
13/10/23 11:12:18 ERROR streaming.StreamJob: Job not successful. Error: NA
13/10/23 11:12:18 INFO streaming.StreamJob: killJob...
Streaming Command Failed!
```

From this output, you can see that your streaming job failed. We intentionally set this up to walk you through a critical part of this process: debugging.

3. Debug the script

To debug the script, open a web browser and go to the job detail page URL given in the command output. In the example above, that URL is: `http://localhost:50030/jobdetails.jsp?jobid=job_201310231802_0004`. Your URL will differ based on your job.

Hadoop job_201402131458_00...

0.0.0.0:50030/jobdetails.jsp?jobid=job_201402131458_0002

Most Visited ▾ Cloudera Cloudera Manager Hue HDFS NameNode Hadoop JobTracker HBase Master Solr

Hadoop job_201402131458_0002 on 0.0.0.0

User: cloudera

Job Name: streamjob927403696501972136.jar

Job File: hdfs://localhost.localdomain:8020/user/cloudera/.staging/job_201402131458_0002/job.xml

Submit Host: localhost.localdomain

Submit Host Address: 127.0.0.1

Job-ACLs: All users are allowed

Job Setup: Successful

Status: Failed

Failure Info: NA

Started at: Thu Feb 13 15:41:57 PST 2014

Failed at: Thu Feb 13 15:43:03 PST 2014

Failed in: 1mins, 5sec

Job Cleanup: Successful

Black-listed TaskTrackers: 1

Kind	% Complete	Num Tasks	Pending	Running	Complete	Killed	Failed/Killed Task Attempts
map	100.00%	20	0	0	0	20	7 / 1
reduce	100.00%	1	0	0	0	1	0 / 0

	Counter	Map	Reduce	Total
Job Counters	Failed map tasks	0	0	1
	Launched map tasks	0	0	8
	Data-local map tasks	0	0	8
	Total time spent by all maps in occupied slots (ms)	0	0	112,654
	Total time spent by all reduces in occupied slots (ms)	0	0	0
	Total time spent by all maps waiting after reserving slots (ms)	0	0	0
	Total time spent by all reduces waiting after reserving slots (ms)	0	0	0

image alt text

From the job detail page, you can see that all 20 map tasks were killed after 8 failed attempts. (The last failed attempt may also be killed, so it may be counted as 7 failed and 1 killed.) If you click on the link for the failed attempts, you get a job failures page where you can see from the stack trace in the error column that the problem was a broken pipe.

Hadoop job_201402131458_0002 failures on 0.0.0.0

Attempt	Task	Machine	State	Error	Logs
attempt_201402131458_0002_m_000000_0	task_201402131458_0002_m_000000	localhost.localdomain	FAILED	java.io.IOException: log:null R/N/S-347/B/B in:NA [rec/s] out:NA [rec/s] writeWrittenToEnabledSkip_0223372036854775807 LOGNAME=null HOST=null USER=null HADOOP_USER=null Last Hadoop input: [null] Last tool output: [null] Date: Thu Feb 13 15:42:18 PST 2014 java.io.IOException: Broken pipe at java.io.FileOutputStream.writeBytes(Native Method) at java.io.FileOutputStream.write(FileOutputStream.java:282) at java.io.BufferedOutputStream.write(BufferedOutputStream.java:185) at java.io.BufferedOutputStream.flushBuffer(BufferedOutputStream.java:65) at java.io.BufferedOutputStream.write(BufferedOutputStream.java:189) at java.io.DataOutputStream.write(DataOutputStream.java:90) at org.apache.hadoop.streaming.io.TextInputWriter.writeUTF8(TextInputWriter.java:72) at org.apache.hadoop.streaming.io.TextInputWriter.writeUTF8(TextInputWriter.java:51) at org.apache.hadoop.streaming.PipeMapper.map(PipeMapper.java:110) at org.apache.hadoop.mapred.Mapper.run(Mapper.java:58) at org.apache.hadoop.streaming.PipeMap	Last 4KB Last 8KB All
attempt_201402131458_0002_m_000000_1	task_201402131458_0002_m_000000	localhost.localdomain	FAILED	java.io.IOException: log:null R/N/S-347/B/B in:NA [rec/s] out:NA [rec/s] writeWrittenToEnabledSkip_0223372036854775807 LOGNAME=null HOST=null USER=null HADOOP_USER=null Last Hadoop input: [null] Last tool output: [null] Date: Thu Feb 13 15:42:31 PST 2014 java.io.IOException: Broken pipe at java.io.FileOutputStream.writeBytes(Native Method) at java.io.FileOutputStream.write(FileOutputStream.java:282) at java.io.BufferedOutputStream.write(BufferedOutputStream.java:185) at java.io.BufferedOutputStream.flushBuffer(BufferedOutputStream.java:65) at java.io.BufferedOutputStream.write(BufferedOutputStream.java:189) at java.io.DataOutputStream.write(DataOutputStream.java:90) at org.apache.hadoop.streaming.io.TextInputWriter.writeUTF8(TextInputWriter.java:72) at org.apache.hadoop.streaming.io.TextInputWriter.writeUTF8(TextInputWriter.java:51) at org.apache.hadoop.streaming.PipeMapper.map(PipeMapper.java:110) at org.apache.hadoop.mapred.Mapper.run(Mapper.java:58) at org.apache.hadoop.streaming.PipeMap	Last 4KB Last 8KB All

image alt text

A broken pipe in a Hadoop streaming job means there was a problem with the script. If you click on one of the links for the task logs in the last column, you finally get to the root of the issue.

Task Logs: 'attempt_201402131458_0002_m_000000_0'

stdout logs

stderr logs

```
Traceback (most recent call last):
  File "/mapred/local/taskTracker/cloudera/jobcache/job_201402131458_0002/attempt_201402131458_0002_m_000000_0/work/./summary_map.py", line 9, in <module>
    data = json.loads(line)
  File "/usr/lib64/python2.6/json/_init_.py", line 307, in loads
    return default_decoder.decode(s)
  File "/usr/lib64/python2.6/json/decoder.py", line 319, in decode
    obj, end = self.raw_decode(s, idx=_w(s, 0).end())
  File "/usr/lib64/python2.6/json/decoder.py", line 336, in raw_decode
    obj, end = self._scanner.iterscan(s, **kw).next()
  File "/usr/lib64/python2.6/json/scanner.py", line 55, in iterscan
    rval, next_pos = action(m, context)
  File "/usr/lib64/python2.6/json/decoder.py", line 183, in JSONObject
    value, end = iterscan(s, idx=end, context=context).next()
  File "/usr/lib64/python2.6/json/scanner.py", line 55, in iterscan
    rval, next_pos = action(m, context)
  File "/usr/lib64/python2.6/json/decoder.py", line 180, in JSONObject
    raise ValueError(errmsgf("Expecting: delimiter", s, end))
ValueError: Expecting: delimiter: line 1 column 91 (char 91)
```

image alt text

In the `stderr` section, note that the `json.loads()` function is throwing a `ValueError` that reads the JSON parser was expecting a colon but received something else. This sounds like there may be errors in the JSON in the logs.

To locate the error(s), you should first get an example line that is broken.

4. Fix the script

Open the `summary_map.py` file and modify it as follows (changes in bold red):

```
#!/usr/bin/python

import json
import sys

# Read all lines from stdin
for line in sys.stdin:
    try:
        # Parse the JSON
        data = json.loads(line)

        # Emit every field
        for field in data.keys():
            print field
    except ValueError:
        # Log the error so we can see it
        sys.stderr.write("%s\n" % line)
        exit(1)
```

The call to `exit()` prevents the job from writing the entire data set into the log files if every line has an error.

5. Rerun the job

Delete the Hadoop output directory and rerun the job with the following commands:

```
$ hadoop fs -rm -R summary
Moved: 'hdfs://localhost.localdomain:8020/user/cloudera/summary' to trash at:
hdfs://localhost.localdomain:8020/user/cloudera/.Trash/Current
$ hadoop jar $STREAMING -input data/heckle/ -input data/jeckle/ -output summary -mapper summary_map.py
  -file summary_map.py -reducer "uniq -c"
...
```

The output should be the same as before. If you follow the same procedure as before to access the logs, you'll now find an example broken line in the `stderr` section.

6. Isolate the broken line

Copy that line and paste it into a file titled `bad.json` which you can save in the current directory. Now that you have an example, run it through the Python JSON formatter, a convenient tool included with Python that will print out JSON formatted input in an easily human-readable format:


```
$ python -mjson.tool < bad.json
Expecting : delimiter: line 1 column 91 (char 91)
```

As expected, it reports the same error as your MapReduce job. Open `bad.json` in an editor and go to column 91 of the first line. At column 91, you'll see an extra double quote around the `itemId` field. You can now test if removing the double-quote will fix the JSON. The `sed` command is an inline editor that you can use to easily remove the duplicate quotes.

Remove the duplicate quotes with the following `sed` command:

```
$ sed 's/"/"/' bad.json | python -mjson.tool
{
  "auth": "15a63c4:e66189ba",
  "createdAt": "2013-05-12T00:00:01-08:00",
  "payload": {
    "itemId": "15607",
    "marker": 240
  },
  "refId": "47c7e2f6",
  "sessionId": "82ada851-0b3c-4e9d-b8cf-0f0a2ebed278",
  "type": "Play",
  "user": 22700996,
  "userAgent": "Mozilla/4.0 (compatible; MSIE 7.0; Windows NT 5.1; .NET CLR 1.1.4322; .NET CLR 2.0.50727; InfoPath.1)"
}
```

It would appear that the double-quote is the only error in that line.

You now have two choices. You can either run a MapReduce job to remove the double quote from all lines where it appears in the data set, or you can correct for the error in-line as you go. Either approach is fine. These next commands show you how to fix it inline.

7. Fix the script

Open up the `summary_map.py` file again and modify it as follows (changes in bold red):

```
#!/usr/bin/python
import json
import sys
# Read all lines from stdin
for line in sys.stdin:
    try:
        # Parse the JSON after fixing the quotes
        data = json.loads(line.replace('""', ''))
        # Emit every field
        for field in data.keys():
            print field
    except ValueError:
        # Log the error so we can see it
        sys.stderr.write("%s\n" % line)
        exit(1)
```

8. Rerun the job

Delete the output directory and rerun the job with the following commands:

```
$ hadoop fs -rm -R summary
Moved: 'hdfs://localhost.localdomain:8020/user/cloudera/summary' to trash at:
hdfs://localhost.localdomain:8020/user/cloudera/.Trash/Current

$ hadoop jar $STREAMING -input data/heckle/ -input data/jeckle/ -output summary -mapper summary_map.py
-file summary_map.py -reducer "uniq -c"
packageJobJar: [../summary_map.py, /tmp/hadoop-training/hadoop-unjar5871949600019015687/] []
/var/folders/1l/x167db9x2rs47q5fs4b255rr0000gp/T/streamjob6407431169321566655.jar tmpDir=null
```

```

13/10/23 13:23:56 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments.
Applications should implement Tool for the same.
13/10/23 13:23:56 INFO mapred.FileInputFormat: Total input paths to process : 20
13/10/23 13:23:56 INFO streaming.StreamJob: getLocalDirs(): [/tmp/hadoop-training/mapred/local]
13/10/23 13:23:56 INFO streaming.StreamJob: Running job: job_201310231802_0006
13/10/23 13:23:56 INFO streaming.StreamJob: To kill this job, run:
13/10/23 13:23:56 INFO streaming.StreamJob: /usr/bin/hadoop job -
Dmapred.job.tracker=192.168.56.101:8021 -kill job_201310231802_0006
13/10/23 13:23:56 INFO streaming.StreamJob: Tracking URL: http://localhost:50030/jobdetails.jsp?
jobid=job_201310231802_0006
13/10/23 13:23:57 INFO streaming.StreamJob: map 0% reduce 0%
13/10/23 13:24:07 INFO streaming.StreamJob: map 3% reduce 0%
13/10/23 13:24:10 INFO streaming.StreamJob: map 5% reduce 0%
13/10/23 13:24:13 INFO streaming.StreamJob: map 6% reduce 0%
13/10/23 13:24:16 INFO streaming.StreamJob: map 8% reduce 0%
13/10/23 13:24:19 INFO streaming.StreamJob: map 9% reduce 0%
13/10/23 13:24:21 INFO streaming.StreamJob: map 10% reduce 0%
...
13/10/23 13:26:18 INFO streaming.StreamJob: map 80% reduce 23%
13/10/23 13:26:22 INFO streaming.StreamJob: map 90% reduce 23%
13/10/23 13:26:24 INFO streaming.StreamJob: map 90% reduce 30%
13/10/23 13:26:25 INFO streaming.StreamJob: map 100% reduce 30%
13/10/23 13:26:27 INFO streaming.StreamJob: map 100% reduce 33%
13/10/23 13:26:30 INFO streaming.StreamJob: map 100% reduce 74%
13/10/23 13:26:33 INFO streaming.StreamJob: map 100% reduce 91%
13/10/23 13:26:35 INFO streaming.StreamJob: map 100% reduce 100%
13/10/23 13:26:37 INFO streaming.StreamJob: Job complete: job_201310231802_0006
13/10/23 13:26:37 INFO streaming.StreamJob: Output: summary

```

To view the results, cat them from HDFS:

```

$ hadoop fs -cat summary/part\*
351712 auth
157151 craetedAt
194561 createdAt
261161 created_at
609352 payload
351712 refId
351712 sessionID
261161 session_id
612873 type
612873 user
351712 userAgent
261161 user_agent

```

There are several things you should notice in this output. First, the type and user fields are the only fields that appear in all lines. Second, you can see that there are three different creation timestamp fields, two different session id fields, and two different user agent fields. Further, some log lines use camel case (e.g., userAgent), and some use snake case (e.g., user_agent). One field name contains a typo that was later corrected. If you add those groups of fields together, you get the total number of lines, so every line has some version of a creation timestamp field, some version of a session id field, and some version of a user agent field. Third, the payload, authorization, and reference id fields are **not** present in every line. From your earlier exploration, you know that the payload field contains additional type-dependent information. The authorization and reference id fields both show up the same number of times, which is also the number of times the camel case user agent field appears.

To get a better feel for the various data types, you can extend your summary job.

9. Extend your summary job

Open the `summary_map.py` file and add the lines as follows (changes in bold red):

```

#!/usr/bin/python
import json
import sys
# Read all lines from stdin

```

```

for line in sys.stdin:
    try:
        # Parse the JSON after fixing the quotes
        data = json.loads(line.replace('""', ''))
        for field in data.keys():
            if field == 'type':
                # Just emit the type field when we see it
                print "%s" % (data[field])
            else:
                # Normalize the file name
                real = field
                if real == 'user_agent':
                    real = 'userAgent'
                elif real == 'session_id':
                    real = 'sessionID'
                elif real == 'created_at' or real == 'craetedAt':
                    real = 'createdAt'
                # Emit the normalized field
                print "%s:%s" % (data['type'], real)
                # Emit all subfields, if there are any
                if type(data[field]) is dict:
                    for subfield in data[field]:
                        print "%s:%s:%s" % (data['type'], real, subfield)
    except ValueError:
        # Log the error so we can see it
        sys.stderr.write("%s\n" % line)
        exit(1)

```

There are two changes here. The first normalizes the field names. The second prints out all fields and subfields for each event type.

10. Run the job

Once you run this job, you should have a pretty good idea of the structure of the log files. Run the job with the following commands:

```

$ hadoop jar $STREAMING -input data/heckle/ -input data/jeckle/ -output summary2 -mapper
summary_map.py -file summary_map.py -reducer "uniq -c"

packageJobJar: [../summary_map.py, /tmp/hadoop-training/hadoop-unjar5795014006503167150/] []
/var/folders/1l/xl67db9x2rs47q5fs4b255rr0000gp/T/streamjob3395742458965177272.jar tmpDir=null
13/10/23 15:11:56 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments.
Applications should implement Tool for the same.
13/10/23 15:11:56 INFO mapred.FileInputFormat: Total input paths to process : 20
13/10/23 15:11:56 INFO streaming.StreamJob: getLocalDirs(): [/tmp/hadoop-training/mapred/local]
13/10/23 15:11:56 INFO streaming.StreamJob: Running job: job_201310231802_0007
13/10/23 15:11:56 INFO streaming.StreamJob: To kill this job, run:
13/10/23 15:11:56 INFO streaming.StreamJob: /usr/bin/hadoop job -Dmapred.job.tracker=localhost:8021 -
kill job_201310231802_0007
13/10/23 15:11:56 INFO streaming.StreamJob: Tracking URL: http://localhost:50030/jobdetails.jsp?
jobid=job_201310231802_0007
13/10/23 15:11:57 INFO streaming.StreamJob: map 0% reduce 0%
13/10/23 15:12:07 INFO streaming.StreamJob: map 3% reduce 0%
13/10/23 15:12:10 INFO streaming.StreamJob: map 4% reduce 0%
...
13/10/23 15:14:41 INFO streaming.StreamJob: map 100% reduce 33%
13/10/23 15:14:44 INFO streaming.StreamJob: map 100% reduce 71%
13/10/23 15:14:47 INFO streaming.StreamJob: map 100% reduce 87%
13/10/23 15:14:51 INFO streaming.StreamJob: map 100% reduce 100%
13/10/23 15:14:52 INFO streaming.StreamJob: Job complete: job_201310231802_0007
13/10/23 15:14:52 INFO streaming.StreamJob: Output: summary2

```

You can then read the files from HDFS and output their contents to standard output. Run the following command:

```
$ hadoop fs -cat summary2/part\  
177 Account  
99 Account:auth  
177 Account:createdAt  
177 Account:payload  
134 Account:payload:new  
134 Account:payload:old  
177 Account:payload:subAction  
99 Account:refId  
177 Account:sessionId  
177 Account:user  
177 Account:userAgent  
5091 AddToQueue  
2919 AddToQueue:auth  
5091 AddToQueue:createdAt  
5091 AddToQueue:payload  
2919 AddToQueue:payload:itemId  
2172 AddToQueue:payload:item_id  
2919 AddToQueue:refId  
5091 AddToQueue:sessionId  
5091 AddToQueue:user  
5091 AddToQueue:userAgent  
3062 Advance  
3062 Advance:createdAt  
3062 Advance:payload  
3062 Advance:payload:item_id  
3062 Advance:payload:marker  
3062 Advance:sessionId  
3062 Advance:user  
3062 Advance:userAgent  
5425 Home  
3109 Home:auth  
5425 Home:createdAt  
5425 Home:payload  
5425 Home:payload:popular  
5425 Home:payload:recent  
5425 Home:payload:recommended  
3109 Home:refId  
5425 Home:sessionId  
5425 Home:user  
5425 Home:userAgent  
19617 Hover  
11376 Hover:auth  
19617 Hover:createdAt  
19617 Hover:payload  
11376 Hover:payload:itemId  
8241 Hover:payload:item_id  
11376 Hover:refId  
19617 Hover:sessionId  
19617 Hover:user  
19617 Hover:userAgent  
274 ItemPage  
154 ItemPage:auth  
274 ItemPage:createdAt  
274 ItemPage:payload  
154 ItemPage:payload:itemId  
120 ItemPage:payload:item_id  
154 ItemPage:refId  
274 ItemPage:sessionId  
274 ItemPage:user  
274 ItemPage:userAgent  
1057 Login  
603 Login:auth  
1057 Login:createdAt  
603 Login:refId
```

```
1057 Login:sessionID
1057 Login:user
1057 Login:userAgent
1018 Logout
  571 Logout:auth
1018 Logout:createdAt
  571 Logout:refId
1018 Logout:sessionID
1018 Logout:user
1018 Logout:userAgent
4424 Pause
2543 Pause:auth
4424 Pause:createdAt
4424 Pause:payload
2543 Pause:payload:itemId
1881 Pause:payload:item_id
4424 Pause:payload:marker
2543 Pause:refId
4424 Pause:sessionID
4424 Pause:user
4424 Pause:userAgent
558568 Play
323244 Play:auth
558568 Play:createdAt
558568 Play:payload
307805 Play:payload:itemId
235324 Play:payload:item_id
543129 Play:payload:marker
323244 Play:refId
558568 Play:sessionID
558568 Play:user
558568 Play:userAgent
  164 Position
  164 Position:createdAt
  164 Position:payload
  164 Position:payload:item_id
  164 Position:payload:marker
  164 Position:sessionId
  164 Position:user
  164 Position:userAgent
1313 Queue
  735 Queue:auth
1313 Queue:createdAt
  735 Queue:refId
1313 Queue:sessionID
1313 Queue:user
1313 Queue:userAgent
652 Rate
  387 Rate:auth
652 Rate:createdAt
652 Rate:payload
  387 Rate:payload:itemId
265 Rate:payload:item_id
652 Rate:payload:rating
  387 Rate:refId
652 Rate:sessionID
652 Rate:user
652 Rate:userAgent
1344 Recommendations
  784 Recommendations:auth
1344 Recommendations:createdAt
1344 Recommendations:payload
1344 Recommendations:payload:recs
  784 Recommendations:refId
1344 Recommendations:sessionID
1344 Recommendations:user
```

```

1344 Recommendations:userAgent
1774 Resume
1774 Resume:createdAt
1774 Resume:payload
1774 Resume:payload:item_id
1774 Resume:payload:marker
1774 Resume:sessionId
1774 Resume:user
1774 Resume:userAgent
1328 Search
769 Search:auth
1328 Search:createdAt
1328 Search:payload
1328 Search:payload:results
769 Search:refId
1328 Search:sessionId
1328 Search:user
1328 Search:userAgent
7178 Stop
4187 Stop:auth
7178 Stop:createdAt
7178 Stop:payload
4187 Stop:payload:itemId
2991 Stop:payload:item_id
7178 Stop:payload:marker
4187 Stop:refId
7178 Stop:sessionId
7178 Stop:user
7178 Stop:userAgent
133 VerifyPassword
78 VerifyPassword:auth
133 VerifyPassword:createdAt
78 VerifyPassword:refId
133 VerifyPassword:sessionId
133 VerifyPassword:user
133 VerifyPassword:userAgent
274 WriteReview
154 WriteReview:auth
274 WriteReview:createdAt
274 WriteReview:payload
154 WriteReview:payload:itemId
120 WriteReview:payload:item_id
274 WriteReview:payload:length
274 WriteReview:payload:rating
154 WriteReview:refId
274 WriteReview:sessionId
274 WriteReview:user
274 WriteReview:userAgent

```

In this output you find some answers to your previous questions and see that new questions arise. First, you can see that the payload subfields also have the camel case/snake case issue. Second, the payload field is only missing in the "Queue," "Login," "Logout," and "VerifyPassword" events, where it is completely absent. Third, the reference id and authorization fields are missing in some records of every event type. Interestingly, they are present exactly as often as the camel case fields. It would appear that they were either added or removed when the case switch was made.

Just to be thorough, you'll need to go one level deeper on your data profile and add some summary statistics on the various job fields. While you're at it, you'll also modify the summary job to normalize the payload subfields before you run it one more time.

11. Normalize the subfields and add the field and subfield values to your output

First, the edit the `summary_map.py` script again. Open the script and add the lines as follows (changes in bold red):

```

#!/usr/bin/python
import json
import sys

```

```

# Read all lines from stdin
for line in sys.stdin:
    try:
        # Parse the JSON after fixing the quotes
        data = json.loads(line.replace('""', ''))
        for field in data.keys():
            if field == 'type':
                # Just emit the type field when we see it
                print "%s" % (data[field])
            else:
                # Normalize the file name
                real = field
                if real == 'user_agent':
                    real = 'userAgent'
                elif real == 'session_id':
                    real = 'sessionID'
                elif real == 'created_at' or real == 'craetedAt':
                    real = 'createdAt'

                # Emit all subfields, if there are any
                if type(data[field]) is dict:
                    print "%s:%s" % (data['type'], real)
                    # Normalize and print the subfields
                    for subfield in data[field]:
                        subreal = subfield
                        if subreal == 'item_id':
                            subreal = 'itemId'

                        print "%s:%s:%s\t%s" % (data['type'], real, subreal, data[field][subfield])
                    else:
                        # Emit the normalized field
                        print "%s:%s\t%s" % (data['type'], real, data[field])

    except ValueError:
        # Log the error so we can see it
        sys.stderr.write("%s\n" % line)
        exit(1)

```

The changes to the script accomplish two tasks: first, it normalizes the subfields; second, it adds the field and subfield values to the output.

12. Write a reduce script

Because you want to calculate summary statistics when you process the output, you'll also need a reduce script. Your reduce script should do all of the same counting as before; additionally, it should summarize the field and subfield values. It should track the values for each field in the data and attempt to identify whether the field is a date, number, category, or identifier. Categorical fields have a small number of possible values, and identifiers have no discernable value pattern.

Create a file called `summary_reduce.py` and enter the following contents. Read the code comments to understand each section of the script.

```

#!/usr/bin/python
import dateutil.parser
import re
import sys
"""
Figure out what type the field is and print appropriate summary stats.
"""
def print_summary():
    if is_heading:
        print "%s - %d" % (last, count)
    elif is_date:
        print "%s - min: %s, max %s, count: %d" % (last, min, max, count)
    elif is_number:
        print "%s - min: %d, max %d, average: %.2f, count: %d" % (last, min, max, float(sum)/count, count)

```

```

elif is_value:
    print "%s - %s, count: %d" % (last, list(values), count)
else:
    print "%s - identifier, count: %d" % (last, count)
last = None
values = set()
is_date = True
is_number = False
is_value = False
is_heading = True
min = None
max = None
sum = 0
count = 0
# Pattern to match a date time field
date_pattern = re.compile(r'\d{4}-\d{2}-\d{2}T\d{2}:\d{2}:\d{2}(Z|[-+]\d{2}:\d{2})')
# Read all lines from stdin
for line in sys.stdin:
    # Split on tab. The first part is the key. The second is the value.
    parts = line.strip().split('\t')
    if parts[0] != last:
        # If there's a previous key, print its summary
        if last != None:
            print_summary()
        # Reset all the summary stats variables
        last = parts[0]
        values = set()
        is_date = True
        is_number = False
        is_identifier = False
        is_heading = True
        min = None
        max = None
        sum = 0
        count = 0
    # Increment the number of times we've seen the field
    count += 1
    # If there was a value of non-zero length, process it
    if len(parts) > 1 and len(parts[1]) > 0:
        is_heading = False
        # If we think it's a date, test if this value parses as a date
        if is_date:
            if date_pattern.match(parts[1]):
                try:
                    tstamp = dateutil.parser.parse(parts[1])
                    # If it does, update the summary stats.
                    if min == None or tstamp < min:
                        min = tstamp
                    if max == None or tstamp > max:
                        max = tstamp
                except (TypeError, ValueError):
                    # If it doesn't parse, then assume it's a number
                    is_date = False
                    is_number = True
                    min = None
                    max = None
            else:
                # If it doesn't match, then assume it's a number
                is_date = False
                is_number = True
                min = None
                max = None

        # If we think it's a number, test if this value parses as a number
        if is_number:
            try:

```



```

num = int(parts[1])
sum += num
# If so, update the summary stats
if min == None or num < min:
    min = num
if max == None or num > max:
    max = num
except ValueError:
    # If not, assume it's categorical
    is_number = False
    is_value = True
# If we think it's categorical, and this value to the category set
if is_value:
    values.add(parts[1])
    # If there are too many categories, call it an identifier
    if len(values) > 10:
        is_value = False
        values = None
# Print the summary for the last key
print_summary()

```

This reduce script tracks every field and the sorts of the field's values. It attempts to identify whether a field is numeric, categorical, or is an identifier. For date fields it tracks the begin and end dates of the range. For numeric fields, it computes the minimum and maximum values and the average. For categorical fields it lists the possible values. For all fields it lists the incidence count.

13. Run the job

Run the new job with the following commands:

```

$ hadoop jar $STREAMING -input data/heckle/ -input data/jeckle/ -output summary3 -mapper
summary_map.py -file summary_map.py -reducer summary_reduce.py -file summary_reduce.py
...
$ hadoop fs -cat summary3/part\*
Account - 177
Account:auth - identifier, count: 99
Account:createdAt - min: 2013-05-06 08:02:56+00:00, max 2013-05-12 23:06:11-08:00, count: 177
Account:payload - 177
Account:payload:new - ['kid'], count: 134
Account:payload:old - ['adult'], count: 134
Account:payload:subAction - ['updatePassword', 'parentalControls', 'updatePaymentInfo'], count: 177
Account:refId - identifier, count: 99
Account:sessionId - identifier, count: 177
Account:user - min: 1634306, max 99314647, average: 47623573.86, count: 177
Account:userAgent - identifier, count: 177
AddToQueue - 5091
AddToQueue:auth - identifier, count: 2919
AddToQueue:createdAt - min: 2013-05-06 08:00:32+00:00, max 2013-05-12 23:57:46-08:00, count: 5091
AddToQueue:payload - 5091
AddToQueue:payload:itemId - identifier, count: 5091
AddToQueue:refId - identifier, count: 2919
AddToQueue:sessionId - identifier, count: 5091
AddToQueue:user - min: 1169676, max 99985450, average: 50609784.80, count: 5091
AddToQueue:userAgent - identifier, count: 5091
Advance - 3062
Advance:createdAt - min: 2013-05-06 08:02:10+00:00, max 2013-05-09 07:27:37+00:00, count: 3062
Advance:payload - 3062
Advance:payload:itemId - identifier, count: 3062
Advance:payload:marker - min: 0, max 8491, average: 2085.52, count: 3062
Advance:sessionId - identifier, count: 3062
Advance:user - min: 1091145, max 99856025, average: 51177856.67, count: 3062
Advance:userAgent - identifier, count: 3062
Home - 5425
Home:auth - identifier, count: 3109
Home:createdAt - min: 2013-05-06 08:00:08+00:00, max 2013-05-12 23:57:18-08:00, count: 5425
Home:payload - 5425

```

Home:payload:popular - min: 1094, max 39475, average: 21398.60, count: 27125
Home:payload:recent - identifier, count: 5222
Home:payload:recommended - identifier, count: 27125
Home:refId - identifier, count: 3109
Home:sessionId - identifier, count: 5425
Home:user - min: 1091145, max 99985450, average: 49984786.82, count: 5425
Home:userAgent - identifier, count: 5425
Hover - 19617
Hover:auth - identifier, count: 11376
Hover:createdAt - min: 2013-05-06 08:00:18+00:00, max 2013-05-12 23:58:08-08:00, count: 19617
Hover:payload - 19617
Hover:payload:itemId - identifier, count: 19617
Hover:refId - identifier, count: 11376
Hover:sessionId - identifier, count: 19617
Hover:user - min: 1091145, max 99985450, average: 50158191.38, count: 19617
Hover:userAgent - identifier, count: 19617
ItemPage - 274
ItemPage:auth - identifier, count: 154
ItemPage:createdAt - min: 2013-05-06 08:02:15+00:00, max 2013-05-12 23:34:12-08:00, count: 274
ItemPage:payload - 274
ItemPage:payload:itemId - identifier, count: 274
ItemPage:refId - identifier, count: 154
ItemPage:sessionId - identifier, count: 274
ItemPage:user - min: 1263067, max 99270605, average: 49770339.91, count: 274
ItemPage:userAgent - identifier, count: 274
Login - 1057
Login:auth - identifier, count: 603
Login:createdAt - min: 2013-05-06 08:01:42+00:00, max 2013-05-12 23:36:07-08:00, count: 1057
Login:refId - identifier, count: 603
Login:sessionId - identifier, count: 1057
Login:user - min: 1091145, max 99856025, average: 49325068.51, count: 1057
Login:userAgent - identifier, count: 1057
Logout - 1018
Logout:auth - identifier, count: 571
Logout:createdAt - min: 2013-05-06 08:13:44+00:00, max 2013-05-12 23:59:15-08:00, count: 1018
Logout:refId - identifier, count: 571
Logout:sessionId - identifier, count: 1018
Logout:user - min: 1091145, max 99985450, average: 48915219.96, count: 1018
Logout:userAgent - identifier, count: 1018
Pause - 4424
Pause:auth - identifier, count: 2543
Pause:createdAt - min: 2013-05-06 08:00:49+00:00, max 2013-05-12 23:57:22-08:00, count: 4424
Pause:payload - 4424
Pause:payload:itemId - identifier, count: 4424
Pause:payload:marker - min: 1, max 7215, average: 2207.71, count: 4424
Pause:refId - identifier, count: 2543
Pause:sessionId - identifier, count: 4424
Pause:user - min: 1091145, max 99985450, average: 50103317.24, count: 4424
Pause:userAgent - identifier, count: 4424
Play - 558568
Play:auth - identifier, count: 323244
Play:createdAt - min: 2013-05-06 08:00:01+00:00, max 2013-05-12 23:59:59-08:00, count: 558568
Play:payload - 558568
Play:payload:itemId - identifier, count: 543129
Play:payload:marker - min: 0, max 8525, average: 2138.50, count: 543129
Play:refId - identifier, count: 323244
Play:sessionId - identifier, count: 558568
Play:user - min: 1091145, max 99985450, average: 50192151.13, count: 558568
Play:userAgent - identifier, count: 558568
Position - 164
Position:createdAt - min: 2013-05-06 08:25:34+00:00, max 2013-05-09 07:02:48+00:00, count: 164
Position:payload - 164
Position:payload:itemId - identifier, count: 164
Position:payload:marker - min: 0, max 6690, average: 2358.18, count: 164
Position:sessionId - identifier, count: 164
Position:user - min: 1091145, max 99413523, average: 49317538.82, count: 164

Position:userAgent - identifier, count: 164
Queue - 1313
Queue:auth - identifier, count: 735
Queue:createdAt - min: 2013-05-06 08:01:21+00:00, max 2013-05-12 23:36:31-08:00, count: 1313
Queue:refId - identifier, count: 735
Queue:sessionId - identifier, count: 1313
Queue:user - min: 1091145, max 99806989, average: 50424708.80, count: 1313
Queue:userAgent - identifier, count: 1313
Rate - 652
Rate:auth - identifier, count: 387
Rate:createdAt - min: 2013-05-06 08:03:32+00:00, max 2013-05-12 23:36:08-08:00, count: 652
Rate:payload - 652
Rate:payload:itemId - identifier, count: 652
Rate:payload:rating - min: 1, max 5, average: 3.54, count: 652
Rate:refId - identifier, count: 387
Rate:sessionId - identifier, count: 652
Rate:user - min: 1091145, max 99314647, average: 49635732.38, count: 652
Rate:userAgent - identifier, count: 652
Recommendations - 1344
Recommendations:auth - identifier, count: 784
Recommendations:createdAt - min: 2013-05-06 08:00:56+00:00, max 2013-05-12 23:58:00-08:00, count: 1344
Recommendations:payload - 1344
Recommendations:payload:recs - identifier, count: 33600
Recommendations:refId - identifier, count: 784
Recommendations:sessionId - identifier, count: 1344
Recommendations:user - min: 1091145, max 99985450, average: 50165065.09, count: 1344
Recommendations:userAgent - identifier, count: 1344
Resume - 1774
Resume:createdAt - min: 2013-05-06 08:02:04+00:00, max 2013-05-09 07:31:48+00:00, count: 1774
Resume:payload - 1774
Resume:payload:itemId - identifier, count: 1774
Resume:payload:marker - min: 0, max 6917, average: 2250.60, count: 1774
Resume:sessionId - identifier, count: 1774
Resume:user - min: 1091145, max 99985450, average: 51027539.16, count: 1774
Resume:userAgent - identifier, count: 1774
Search - 1328
Search:auth - identifier, count: 769
Search:createdAt - min: 2013-05-06 08:02:11+00:00, max 2013-05-12 23:36:56-08:00, count: 1328
Search:payload - 1328
Search:payload:results - identifier, count: 26560
Search:refId - identifier, count: 769
Search:sessionId - identifier, count: 1328
Search:user - min: 1170207, max 99976229, average: 50523812.45, count: 1328
Search:userAgent - identifier, count: 1328
Stop - 7178
Stop:auth - identifier, count: 4187
Stop:createdAt - min: 2013-05-06 08:04:10+00:00, max 2013-05-12 23:55:49-08:00, count: 7178
Stop:payload - 7178
Stop:payload:itemId - identifier, count: 7178
Stop:payload:marker - min: 172, max 8233, average: 2692.93, count: 7178
Stop:refId - identifier, count: 4187
Stop:sessionId - identifier, count: 7178
Stop:user - min: 1091145, max 99976229, average: 49769162.34, count: 7178
Stop:userAgent - identifier, count: 7178
VerifyPassword - 133
VerifyPassword:auth - identifier, count: 78
VerifyPassword:createdAt - min: 2013-05-06 10:02:24+00:00, max 2013-05-12 23:02:33-08:00, count: 133
VerifyPassword:refId - identifier, count: 78
VerifyPassword:sessionId - identifier, count: 133
VerifyPassword:user - min: 1634306, max 99314647, average: 47262951.69, count: 133
VerifyPassword:userAgent - identifier, count: 133
WriteReview - 274
WriteReview:auth - identifier, count: 154
WriteReview:createdAt - min: 2013-05-06 08:11:46+00:00, max 2013-05-12 23:38:58-08:00, count: 274
WriteReview:payload - 274
WriteReview:payload:itemId - identifier, count: 274

```
WriteReview:payload:length - min: 52, max 1192, average: 627.63, count: 274
WriteReview:payload:rating - min: 1, max 5, average: 4.03, count: 274
WriteReview:refId - identifier, count: 154
WriteReview:sessionId - identifier, count: 274
WriteReview:user - min: 1263067, max 99270605, average: 49770339.91, count: 274
WriteReview:userAgent - identifier, count: 274
```

14. Review your data summary

You now have a very useful summary of the data from which you can make some statements.

- user ids are in the rough range of 1000000 to 100000000.
- item ids are not all numeric (otherwise they wouldn't have been marked as an identifier).
- the average rating for Rating events is 3.5 and for WriteReview events is 4.0.
- creation timestamps aren't all in the same time zone. Some are in UTC, and some are in UTC-8. There may be others that aren't exposed. You'll have to take that into account when cleaning the data.
- the Account action has three possible sub-actions: update password, update payment info, and parental controls.
- You can also deduce from the other payload fields in the Account event that the parental control actions recorded are only those that enable controls, but none are recorded that disable parental controls (which is rather odd).
- all of the fields and subfields are now normalized
- in the payload field of Play events, you know that the item id and marker subfields are not present in every record. In fact, they are both missing in the same number of records.

Step 5. Determine distinct users and account action details

There are two more aspects of the data you should explore before you work on a more formal data transformation on which you will base the rest of your work: the number of distinct users and sessions and the account action details, especially parental controls.

1. Get user and session counts

You can reuse the `grep_field.sh` script to get the user and session counts. Enter the following commands:

```
$ hadoop jar $STREAMING -D stream.non.zero.exit.is.failure=false -input data/heckle/ -input
data/heckle/ -output users -mapper "grep_field.sh user" -file grep_field.sh -reducer 'bash -c "uniq |
wc -l"'
...
$ hadoop fs -cat users/part\*
2195
$ hadoop jar $STREAMING -D stream.non.zero.exit.is.failure=false -input data/heckle/ -input
data/heckle/ -output sessions -mapper 'grep_field.sh session(ID|_id)' -file grep_field.sh -reducer
'bash -c "uniq | wc -l"'
...
$ hadoop fs -cat sessions/part\*
5308
```

In these jobs, you'll use a slightly more complicated reducer than before. Hadoop streaming doesn't accept pipes directly, so the command must be wrapped in `bash -c`.

To make sure you understand the parental control event, let's look at an example. Enter the following commands:

```
$ grep -m 1 '"parentalControls"' data/heckle/web.log
{"auth": "319e75f:399b2e61", "createdAt": "2013-05-12T00:01:27-08:00", "payload": {"new": "kid",
"old": "adult", "subAction": "parentalControls"}, "refId": "752c4bc5", "sessionId": "81076d3e-ad42-
4567-a5dd-df3d4a0f3274", "type": "Account", "user": 52029279, "userAgent": "Mozilla/4.0 (compatible;
MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR
3.0.30729; .NET4.0C; .NET4.0E; .NET CLR 1.1.4322; InfoPath.3)"} }
```

This example confirms that the new and old fields do belong to the parental controls sub-action. Now, to confirm that our summary script was right about the new and old field values, you use our `grep_field.sh` script one last time. Enter the following commands:

```
$ hadoop jar $STREAMING -D stream.non.zero.exit.is.failure=false -input data/heckle/ -input
```

```
data/jeckle/ -output new -mapper "grep_field.sh new" -file grep_field.sh -reducer uniq
...
$ hadoop fs -cat new/part\*
kid
```

From these results, you can see that “kid” really is the only value for the new Account payload field in the data set, which means that “adult” must really be the only value for the old field. That sounds odd, but it’s what the data contains.

You now have a pretty good understanding of the data set and can proceed with building a clean data set from which you will work for the rest of the exercises.

Navigation

[Table of Contents](#)

[Solution Kit Introduction](#)

[Project Introduction](#)

[Exploring the Data](#)

[Cleaning the Data \(next\)](#)

[Classifying Users](#)

[Clustering Sessions](#)

[Predicting User Ratings \(Building a Recommender\)](#)

[Conclusion](#)

Products

Cloudera Enterprise

Cloudera Express

Cloudera Manager

CDH

All Downloads

Professional Services

Training

Solutions

Enterprise Solutions

Partner Solutions

Industry Solutions

Partners

Resource Library

Support

About

Hadoop & Big Data

Management Team

Board

Events

Press Center

Careers

Contact Us

Subscription Center

English ▼

Follow us:

Share: